

TOWARDS A SYSTEMATIC BENCHMARKING OF ONTOLOGY-BASED QUERY REWRITING SYSTEMS


José Mora and Óscar Corcho

{jmora, ocorcho}@fi.upm.es


Boadilla - October 10, 2013

INDEX

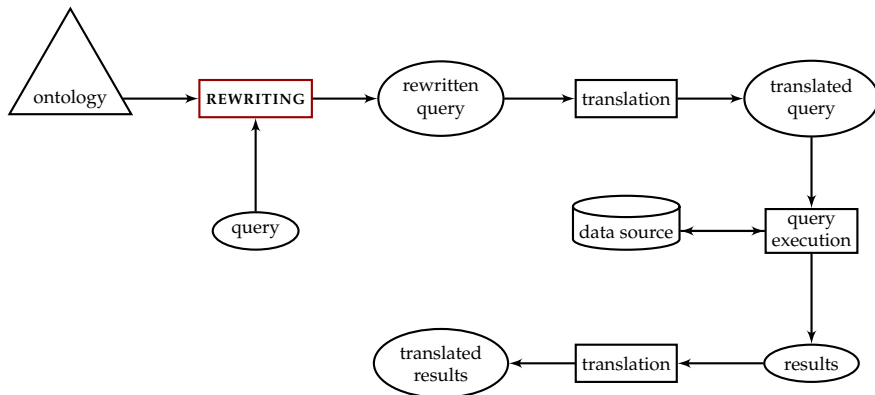
- 1 INTRODUCTION**
- 2 STATE OF THE ART**
- 3 ANALYSIS**
- 4 RESULTS**
- 5 CONCLUSIONS**



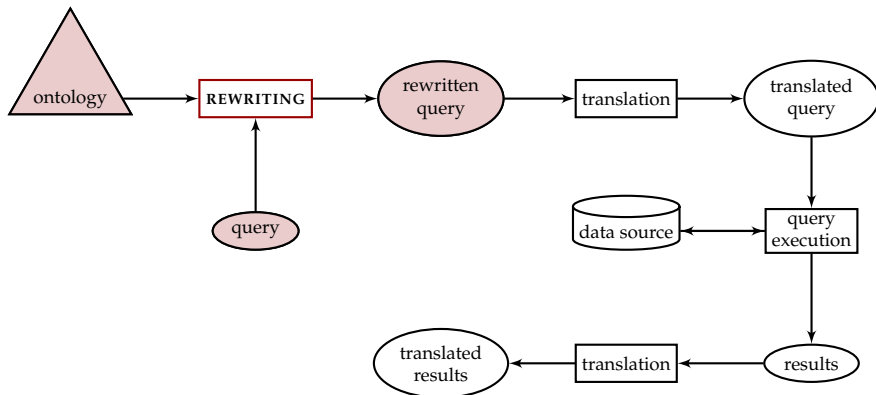
APPROACH



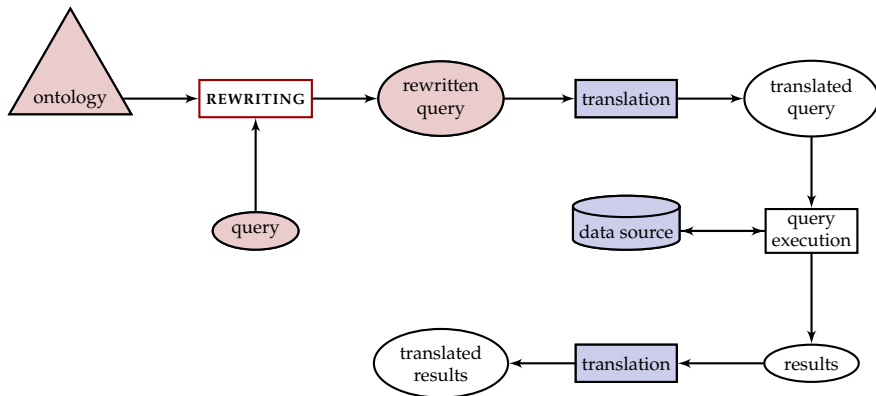
CONTEXT



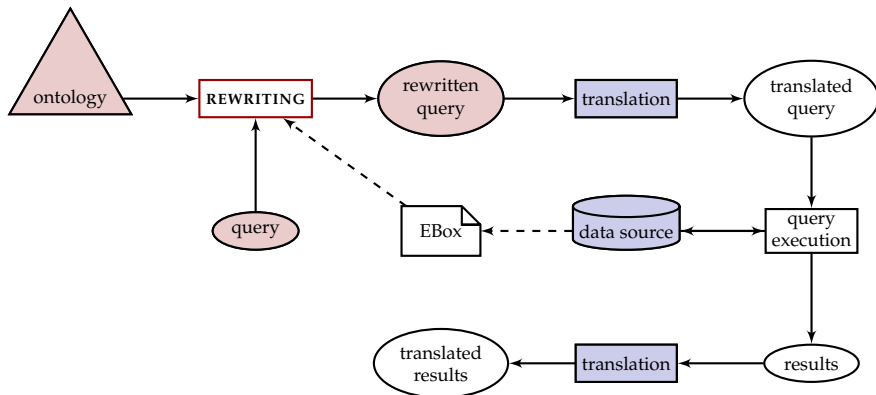
CONTEXT



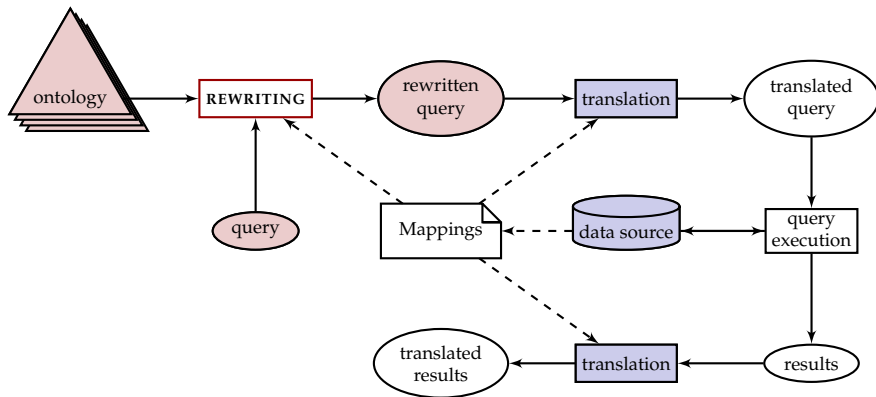
CONTEXT



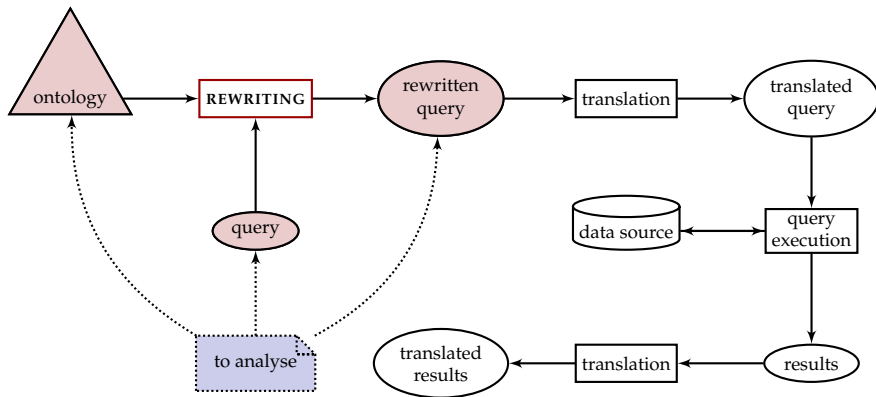
CONTEXT



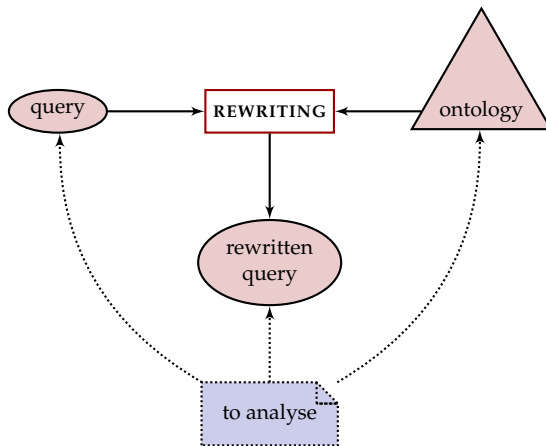
CONTEXT



CONTEXT



FOCUSING



LOGICS

| axiom ¹ \ logic | DL-Lite _{core} | DL-Lite _F | DL-Lite _R | Rapid's | \mathcal{ELHIQ}^- | Datalog ^{±2} | Horn-SHIQ |
|--|-------------------------|----------------------|----------------------|---------|---------------------|-----------------------|-----------|
| $B_1 \sqsubseteq B_2, B_1 \sqsubseteq \neg B_2^3$ | ✓ | ✓ | ✓ | ✓ | ✓ | ✓ | ✓ |
| $\geq 2R_1 \sqsubseteq \perp$ | ✗ | ✓ | ✗ | ✗ | ✗ | ✗ | ✓ |
| $R_1 \sqsubseteq R_2, R_1 \sqsubseteq \neg R_2$ | ✗ | ✗ | ✓ | ✓ | ✓ | ✓ | ✓ |
| $B_1 \sqsubseteq \exists R_1, \exists R_1 \sqsubseteq B_1$ | ✗ | ✗ | ✓ | ✓ | ✓ | ✓ | ✓ |
| $B_1 \sqsubseteq \exists R_1.B_2$ | ✗ | ✗ | ✗ | ✓ | ✓ | ✓ | ✓ |
| $\exists R_1.B_1 \sqsubseteq B_2$ | ✗ | ✗ | ✗ | ✗ | ✓ | ✗ | ✓ |
| $B_1 \sqcap B_2 \sqsubseteq B_3$ | ✗ | ✗ | ✗ | ✗ | ✓ | ✗ | ✓ |
| $\{a\} \sqsubseteq B, B \sqsubseteq \{a\}, B(a)$ | ✗ | ✗ | ✗ | ✗ | ✓ | ✗ | ✓ |
| <i>n</i> -ary predicates | ✗ | ✗ | ✗ | ✗ | ✗ | ✓ | ✗ |
| <i>trans</i> (R_1) | ✗ | ✗ | ✗ | ✗ | ✗ | ✗ | ✓ |
| $B_1 \sqsubseteq \forall R_1.B_2, B_1 \sqsubseteq \leq 1R_1.B_2$ | ✗ | ✗ | ✗ | ✗ | ✗ | ✗ | ✓ |

¹here B_i represents a basic concept and R_j a role that may be basic or inverted.

²as implemented in Nyaya

³note that some systems implementing negation assume a consistent ABox

SYSTEMS

| System | Input | Output | Reference |
|-----------------|-----------------------------------|----------------|-----------------------------|
| Quonto | DL-Lite _R | UCQ | Calvanese et al. (2007) |
| REQUIEM | \mathcal{ELHIQ}^\neg | Datalog or UCQ | Pérez-Urbina et al. (2009) |
| Presto | DL-Lite _R | Datalog | Rosati and Almatelli (2010) |
| Rapid | DL-Lite _R ⁴ | Datalog or UCQ | Chortaras et al. (2011) |
| Nyaya | $Datalog^\pm$ | UCQ | Gottlob et al. (2011) |
| Venetis' | DL-Lite _R | UCQ | Venetis et al. (2011) |
| Prexto | DL-Lite _R and EBox | Datalog or UCQ | Rosati (2012) |
| Clipper | Horn- $SHIQ$ | Datalog | Eiter et al. (2012) |
| kyrie | \mathcal{ELHIQ}^\neg | Datalog or UCQ | Mora and Corcho (2013) |

⁴Close to OWL2 QL, $B_1 \sqsubseteq \exists R.B_2$ axioms are supported

PREVIOUS EFFORTS

- Imprialou et al. (2012) is the most notable previous effort
- Automatic and exhaustive coverage of the ontologies with queries
- Detected flaws in the implementations of most systems
- Still focused on qualitative results (soundness and completeness)

DIMENSIONS



■ Expressiveness in tests: traditionally DL-Lite_R

- Most expressive logic in the intersection of all systems
- Systems that handle more expressive logics cannot show their full potential

■ Expressiveness lost in translation to Datalog

- How expressive are the ontologies for OBDA?
- Consequences of not covered expressiveness on precision and completeness

■ Output complexity: apples, oranges and pears

- How to compare UCQs and Datalog programs?
- Characteristics of the system that is going to execute them

■ Input complexity: treat with care

- What kind of queries can be processed?
- What kind of queries do we want to process?

■ Additional inputs: to each one its own

- EBox, cache of previous queries, etc.
- Comparison among systems and comparison with Reality™

DIMENSIONS

- Expressiveness in tests: traditionally DL-Lite_R
 - Most expressive logic in the intersection of all systems
 - Systems that handle more expressive logics cannot show their full potential
- Expressiveness lost in translation to Datalog
 - How expressive are the ontologies for OBDA?
 - Consequences of not covered expressiveness on precision and completeness
- Output complexity: apples, oranges and pears
 - How to compare UCQs and Datalog programs?
 - Characteristics of the system that is going to execute them
- Input complexity: treat with care
 - What kind of queries can be processed?
 - What kind of queries do we want to process?
- Additional inputs: to each one its own
 - EBox, cache of previous queries, etc.
 - Comparison among systems and comparison with Reality™

DIMENSIONS

- Expressiveness in tests: traditionally DL-Lite_R
 - Most expressive logic in the intersection of all systems
 - Systems that handle more expressive logics cannot show their full potential
- Expressiveness lost in translation to Datalog
 - How expressive are the ontologies for OBDA?
 - Consequences of not covered expressiveness on precision and completeness
- Output complexity: apples, oranges and pears
 - How to compare UCQs and Datalog programs?
 - Characteristics of the system that is going to execute them
- Input complexity: treat with care
 - What kind of queries can be processed?
 - What kind of queries do we want to process?
- Additional inputs: to each one its own
 - EBox, cache of previous queries, etc.
 - Comparison among systems and comparison with Reality™

DIMENSIONS



- Expressiveness in tests: traditionally DL-Lite_R
 - Most expressive logic in the intersection of all systems
 - Systems that handle more expressive logics cannot show their full potential
- Expressiveness lost in translation to Datalog
 - How expressive are the ontologies for OBDA?
 - Consequences of not covered expressiveness on precision and completeness
- Output complexity: apples, oranges and pears
 - How to compare UCQs and Datalog programs?
 - Characteristics of the system that is going to execute them
- Input complexity: treat with care
 - What kind of queries can be processed?
 - What kind of queries do we want to process?
- Additional inputs: to each one its own
 - EBox, cache of previous queries, etc.
 - Comparison among systems and comparison with Reality™

DIMENSIONS

- Expressiveness in tests: traditionally DL-Lite_R
 - Most expressive logic in the intersection of all systems
 - Systems that handle more expressive logics cannot show their full potential
- Expressiveness lost in translation to Datalog
 - How expressive are the ontologies for OBDA?
 - Consequences of not covered expressiveness on precision and completeness
- Output complexity: apples, oranges and pears
 - How to compare UCQs and Datalog programs?
 - Characteristics of the system that is going to execute them
- Input complexity: treat with care
 - What kind of queries can be processed?
 - What kind of queries do we want to process?
- Additional inputs: to each one its own
 - EBox, cache of previous queries, etc.
 - Comparison among systems and comparison with Reality™

DIMENSIONS



- Expressiveness in tests: traditionally DL-Lite_R
 - Most expressive logic in the intersection of all systems
 - Systems that handle more expressive logics cannot show their full potential
- Expressiveness lost in translation to Datalog
 - How expressive are the ontologies for OBDA?
 - Consequences of not covered expressiveness on precision and completeness
- Output complexity: apples, oranges and pears
 - How to compare UCQs and Datalog programs?
 - Characteristics of the system that is going to execute them
- Input complexity: treat with care
 - What kind of queries can be processed?
 - What kind of queries do we want to process?
- Additional inputs: to each one its own
 - EBox, cache of previous queries, etc.
 - Comparison among systems and comparison with Reality™

DIMENSIONS



- Expressiveness in tests: traditionally DL-Lite_R
 - Most expressive logic in the intersection of all systems
 - Systems that handle more expressive logics cannot show their full potential
- Expressiveness lost in translation to Datalog
 - How expressive are the ontologies for OBDA?
 - Consequences of not covered expressiveness on precision and completeness
- Output complexity: apples, oranges and pears
 - How to compare UCQs and Datalog programs?
 - Characteristics of the system that is going to execute them
- Input complexity: treat with care
 - What kind of queries can be processed?
 - What kind of queries do we want to process?
- Additional inputs: to each one its own
 - EBox, cache of previous queries, etc.
 - Comparison among systems and comparison with Reality™

DIMENSIONS



- Expressiveness in tests: traditionally DL-Lite_R
 - Most expressive logic in the intersection of all systems
 - Systems that handle more expressive logics cannot show their full potential
- Expressiveness lost in translation to Datalog
 - How expressive are the ontologies for OBDA?
 - Consequences of not covered expressiveness on precision and completeness
- Output complexity: apples, oranges and pears
 - How to compare UCQs and Datalog programs?
 - Characteristics of the system that is going to execute them
- Input complexity: treat with care
 - What kind of queries can be processed?
 - What kind of queries do we want to process?
- Additional inputs: to each one its own
 - EBox, cache of previous queries, etc.
 - Comparison among systems and comparison with Reality™

DIMENSIONS



- Expressiveness in tests: traditionally DL-Lite_R
 - Most expressive logic in the intersection of all systems
 - Systems that handle more expressive logics cannot show their full potential
- Expressiveness lost in translation to Datalog
 - How expressive are the ontologies for OBDA?
 - Consequences of not covered expressiveness on precision and completeness
- Output complexity: apples, oranges and pears
 - How to compare UCQs and Datalog programs?
 - Characteristics of the system that is going to execute them
- Input complexity: treat with care
 - What kind of queries can be processed?
 - What kind of queries do we want to process?
- Additional inputs: to each one its own
 - EBox, cache of previous queries, etc.
 - Comparison among systems and comparison with Reality™

DIMENSIONS



- Expressiveness in tests: traditionally DL-Lite_R
 - Most expressive logic in the intersection of all systems
 - Systems that handle more expressive logics cannot show their full potential
- Expressiveness lost in translation to Datalog
 - How expressive are the ontologies for OBDA?
 - Consequences of not covered expressiveness on precision and completeness
- Output complexity: apples, oranges and pears
 - How to compare UCQs and Datalog programs?
 - Characteristics of the system that is going to execute them
- Input complexity: treat with care
 - What kind of queries can be processed?
 - What kind of queries do we want to process?
- Additional inputs: to each one its own
 - EBox, cache of previous queries, etc.
 - Comparison among systems and comparison with Reality™

DIMENSIONS



- Expressiveness in tests: traditionally DL-Lite_R
 - Most expressive logic in the intersection of all systems
 - Systems that handle more expressive logics cannot show their full potential
- Expressiveness lost in translation to Datalog
 - How expressive are the ontologies for OBDA?
 - Consequences of not covered expressiveness on precision and completeness
- Output complexity: apples, oranges and pears
 - How to compare UCQs and Datalog programs?
 - Characteristics of the system that is going to execute them
- Input complexity: treat with care
 - What kind of queries can be processed?
 - What kind of queries do we want to process?
- Additional inputs: to each one its own
 - EBox, cache of previous queries, etc.
 - Comparison among systems and comparison with Reality™

DIMENSIONS



- Expressiveness in tests: traditionally DL-Lite_R
 - Most expressive logic in the intersection of all systems
 - Systems that handle more expressive logics cannot show their full potential
- Expressiveness lost in translation to Datalog
 - How expressive are the ontologies for OBDA?
 - Consequences of not covered expressiveness on precision and completeness
- Output complexity: apples, oranges and pears
 - How to compare UCQs and Datalog programs?
 - Characteristics of the system that is going to execute them
- Input complexity: treat with care
 - What kind of queries can be processed?
 - What kind of queries do we want to process?
- Additional inputs: to each one its own
 - EBox, cache of previous queries, etc.
 - Comparison among systems and comparison with Reality™

DIMENSIONS



- Expressiveness in tests: traditionally DL-Lite_R
 - Most expressive logic in the intersection of all systems
 - Systems that handle more expressive logics cannot show their full potential
- Expressiveness lost in translation to Datalog
 - How expressive are the ontologies for OBDA?
 - Consequences of not covered expressiveness on precision and completeness
- Output complexity: apples, oranges and pears
 - How to compare UCQs and Datalog programs?
 - Characteristics of the system that is going to execute them
- Input complexity: treat with care
 - What kind of queries can be processed?
 - What kind of queries do we want to process?
- Additional inputs: to each one its own
 - EBox, cache of previous queries, etc.
 - Comparison among systems and comparison with RealityTM

DIMENSIONS



- Expressiveness in tests: traditionally DL-Lite_R
 - Most expressive logic in the intersection of all systems
 - Systems that handle more expressive logics cannot show their full potential
- Expressiveness lost in translation to Datalog
 - How expressive are the ontologies for OBDA?
 - Consequences of not covered expressiveness on precision and completeness
- Output complexity: apples, oranges and pears
 - How to compare UCQs and Datalog programs?
 - Characteristics of the system that is going to execute them
- Input complexity: treat with care
 - What kind of queries can be processed?
 - What kind of queries do we want to process?
- Additional inputs: to each one its own
 - EBox, cache of previous queries, etc.
 - Comparison among systems and comparison with RealityTM

DIMENSIONS



- Expressiveness in tests: traditionally DL-Lite_R
 - Most expressive logic in the intersection of all systems
 - Systems that handle more expressive logics cannot show their full potential
- Expressiveness lost in translation to Datalog
 - How expressive are the ontologies for OBDA?
 - Consequences of not covered expressiveness on precision and completeness
- Output complexity: apples, oranges and pears
 - How to compare UCQs and Datalog programs?
 - Characteristics of the system that is going to execute them
- Input complexity: treat with care
 - What kind of queries can be processed?
 - What kind of queries do we want to process?
- Additional inputs: to each one its own
 - EBox, cache of previous queries, etc.
 - Comparison among systems and comparison with RealityTM

ASSETS



- Several assets used for evaluation⁵
- Usual ontologies (A, AX, P1, P5, P5X, S, U, UX, V)
- Not so usual ontologies (core, galen-lite)
- New ontologies (AXE, AXEb, P5XE, UXE)
- Usually with 5 queries, but up to 9 in some cases

⁵available here: <http://j.mp/benchmarkqueryrewriting>

ASSETS



- Several assets used for evaluation⁵
- Usual ontologies (A, AX, P1, P5, P5X, S, U, UX, V)
- Not so usual ontologies (core, galen-lite)
- New ontologies (AXE, AXEb, P5XE, UXE)
- Usually with 5 queries, but up to 9 in some cases

⁵available here: <http://j.mp/benchmarkqueryrewriting>

ASSETS



- Several assets used for evaluation⁵
- Usual ontologies (A, AX, P1, P5, P5X, S, U, UX, V)
- Not so usual ontologies (core, galen-lite)
- New ontologies (AXE, AXEb, P5XE, UXE)
- Usually with 5 queries, but up to 9 in some cases

⁵available here: <http://j.mp/benchmarkqueryrewriting>

ASSETS



- Several assets used for evaluation⁵
- Usual ontologies (A, AX, P1, P5, P5X, S, U, UX, V)
- Not so usual ontologies (core, galen-lite)
- New ontologies (AXE, AXEb, P5XE, UXE)
- Usually with 5 queries, but up to 9 in some cases

⁵available here: <http://j.mp/benchmarkqueryrewriting>

ASSETS



- Several assets used for evaluation⁵
- Usual ontologies (A, AX, P1, P5, P5X, S, U, UX, V)
- Not so usual ontologies (core, galen-lite)
- New ontologies (AXE, AXEb, P5XE, UXE)
- Usually with 5 queries, but up to 9 in some cases

⁵available here: <http://j.mp/benchmarkqueryrewriting>

RESULTS I

| \mathcal{O} | q | REQUIEM(F) | | Presto | | Rapid | | Clipper | | kyrie | |
|---------------|-----|------------|------|--------|------|-------|------|---------|------|-------|------|
| | | size | time | size | time | size | time | size | time | size | time |
| U | 1 | 19 | 12 | 4 | 7 | 4 | 3 | 2 | 21 | 2 | 0 |
| | 2 | 47 | 16 | 2 | 9 | 2 | 9 | 49 | 19 | 47 | 3 |
| | 3 | 20 | 9 | 8 | 16 | 8 | 12 | 21 | 24 | 20 | 3 |
| | 4 | 64 | 15 | 3 | 12 | 3 | 3 | 63 | 18 | 64 | 3 |
| | 5 | 53 | 12 | 8 | 15 | 8 | 12 | 53 | 15 | 53 | 0 |
| | 6 | 20 | 12 | 19 | 6 | 21 | 15 | 16 | 18 | 16 | 9 |
| | 7 | 49 | 25 | 22 | 15 | 22 | 18 | 44 | 18 | 45 | 12 |
| | 8 | 10 | 9 | 13 | 6 | 13 | 9 | 10 | 20 | 10 | 3 |
| | 9 | 29 | 15 | 24 | 12 | 24 | 17 | 19 | 20 | 21 | 7 |
| UX | 1 | 22 | 6 | 7 | 10 | 7 | 3 | 5 | 27 | 5 | 6 |
| | 2 | 52 | 15 | 2 | 15 | 2 | 15 | 54 | 27 | 52 | 3 |
| | 3 | 24 | 15 | 10 | 28 | 10 | 9 | 25 | 28 | 24 | 3 |
| | 4 | 70 | 15 | 6 | 19 | 6 | 12 | 69 | 22 | 70 | 0 |
| | 5 | 56 | 15 | 11 | 15 | 11 | 15 | 56 | 21 | 56 | 12 |
| | 6 | 24 | 18 | 28 | 15 | 27 | 22 | 20 | 19 | 20 | 3 |
| | 7 | 55 | 28 | 29 | 21 | 27 | 21 | 50 | 28 | 51 | 12 |
| | 8 | 11 | 6 | 14 | 12 | 14 | 13 | 11 | 26 | 11 | 1 |
| | 9 | 32 | 15 | 30 | 21 | 30 | 15 | 22 | 46 | 24 | 4 |

TABLE : Results of the execution to obtain Datalog (time in ms)

RESULTS I

| \mathcal{O} | q | REQUIEM(F) | | Presto | | Rapid | | Clipper | | kyrie | |
|---------------|-----|------------|------|--------|------|-------|------|---------|------|-------|------|
| | | size | time | size | time | size | time | size | time | size | time |
| U | 1 | 19 | 12 | 4 | 7 | 4 | 3 | 2 | 21 | 2 | 0 |
| | 2 | 47 | 16 | 2 | 9 | 2 | 9 | 49 | 19 | 47 | 3 |
| | 3 | 20 | 9 | 8 | 16 | 8 | 12 | 21 | 24 | 20 | 3 |
| | 4 | 64 | 15 | 3 | 12 | 3 | 3 | 63 | 18 | 64 | 3 |
| | 5 | 53 | 12 | 8 | 15 | 8 | 12 | 53 | 15 | 53 | 0 |
| | 6 | 20 | 12 | 19 | 6 | 21 | 15 | 16 | 18 | 16 | 9 |
| | 7 | 49 | 25 | 22 | 15 | 22 | 18 | 44 | 18 | 45 | 12 |
| | 8 | 10 | 9 | 13 | 6 | 13 | 9 | 10 | 20 | 10 | 3 |
| | 9 | 29 | 15 | 24 | 12 | 24 | 17 | 19 | 20 | 21 | 7 |
| UX | 1 | 22 | 6 | 7 | 10 | 7 | 3 | 5 | 27 | 5 | 6 |
| | 2 | 52 | 15 | 2 | 15 | 2 | 15 | 54 | 27 | 52 | 3 |
| | 3 | 24 | 15 | 10 | 28 | 10 | 9 | 25 | 28 | 24 | 3 |
| | 4 | 70 | 15 | 6 | 19 | 6 | 12 | 69 | 22 | 70 | 0 |
| | 5 | 56 | 15 | 11 | 15 | 11 | 15 | 56 | 21 | 56 | 12 |
| | 6 | 24 | 18 | 28 | 15 | 27 | 22 | 20 | 19 | 20 | 3 |
| | 7 | 55 | 28 | 29 | 21 | 27 | 21 | 50 | 28 | 51 | 12 |
| | 8 | 11 | 6 | 14 | 12 | 14 | 13 | 11 | 26 | 11 | 1 |
| | 9 | 32 | 15 | 30 | 21 | 30 | 15 | 22 | 46 | 24 | 4 |

TABLE : Results of the execution to obtain Datalog (time in ms)

RESULTS II

| \mathcal{O} | q | REQUIEM(F) | | Rapid | | Prexto | | Nyaya | | kyrie | |
|---------------|-----|------------|-------|-------|------|--------|-------|-------|--------|-------|-------|
| | | size | time | size | time | size | time | size | time | size | time |
| U | 1 | 2 | 15 | 2 | 3 | 2 | 9 | 2 | 5 | 2 | 0 |
| | 2 | 1 | 103 | 1 | 15 | 1 | 15 | 1 | 1 | 1 | 34 |
| | 3 | 4 | 212 | 4 | 9 | 4 | 18 | 4 | 34 | 4 | 18 |
| | 4 | 2 | 3762 | 2 | 12 | 2 | 15 | 2 | 4 | 2 | 50 |
| | 5 | 10 | 13034 | 10 | 18 | 10 | 15 | 10 | 33 | 10 | 37 |
| | 6 | 29 | 47 | 29 | 28 | 28 | 12 | 40 | 1595 | 29 | 28 |
| | 7 | 42 | 797 | 42 | 37 | 70 | 18 | 54 | 670 | 42 | 43 |
| | 8 | 10 | 15 | 10 | 18 | 10 | 6 | 10 | 63 | 10 | 3 |
| | 9 | 960 | 1893 | 960 | 209 | 960 | 928 | 960 | 75135 | 960 | 1107 |
| UX | 1 | 5 | 15 | 5 | 12 | 5 | 12 | 5 | 22 | 5 | 9 |
| | 2 | 1 | 172 | 1 | 12 | 1 | 16 | 1 | 3 | 1 | 37 |
| | 3 | 12 | 2062 | 12 | 15 | 12 | 28 | 12 | 55 | 12 | 21 |
| | 4 | 5 | 31422 | 5 | 15 | 5 | 23 | 5 | 6 | 5 | 47 |
| | 5 | 25 | 91878 | 25 | 27 | 25 | 23 | 25 | 39 | 25 | 46 |
| | 6 | 323 | 468 | 323 | 106 | 448 | 178 | 348 | 2685 | 323 | 187 |
| | 7 | 1456 | 37212 | 224 | 81 | 280 | 75 | 264 | 852 | 224 | 121 |
| | 8 | 20 | 21 | 20 | 23 | 20 | 15 | 20 | 61 | 20 | 9 |
| | 9 | 4200 | 30506 | 4200 | 739 | 4200 | 21181 | 4200 | 366673 | 4200 | 16875 |

TABLE : Results of the execution to obtain UCQ (time in ms)

RESULTS II

| \mathcal{O} | q | REQUIEM(F) | | Rapid | | Prexto | | Nyaya | | kyrie | |
|---------------|-----|------------|-------|-------|------|--------|-------|-------|--------|-------|-------|
| | | size | time | size | time | size | time | size | time | size | time |
| U | 1 | 2 | 15 | 2 | 3 | 2 | 9 | 2 | 5 | 2 | 0 |
| | 2 | 1 | 103 | 1 | 15 | 1 | 15 | 1 | 1 | 1 | 34 |
| | 3 | 4 | 212 | 4 | 9 | 4 | 18 | 4 | 34 | 4 | 18 |
| | 4 | 2 | 3762 | 2 | 12 | 2 | 15 | 2 | 4 | 2 | 50 |
| | 5 | 10 | 13034 | 10 | 18 | 10 | 15 | 10 | 33 | 10 | 37 |
| | 6 | 29 | 47 | 29 | 28 | 28 | 12 | 40 | 1595 | 29 | 28 |
| | 7 | 42 | 797 | 42 | 37 | 70 | 18 | 54 | 670 | 42 | 43 |
| | 8 | 10 | 15 | 10 | 18 | 10 | 6 | 10 | 63 | 10 | 3 |
| | 9 | 960 | 1893 | 960 | 209 | 960 | 928 | 960 | 75135 | 960 | 1107 |
| UX | 1 | 5 | 15 | 5 | 12 | 5 | 12 | 5 | 22 | 5 | 9 |
| | 2 | 1 | 172 | 1 | 12 | 1 | 16 | 1 | 3 | 1 | 37 |
| | 3 | 12 | 2062 | 12 | 15 | 12 | 28 | 12 | 55 | 12 | 21 |
| | 4 | 5 | 31422 | 5 | 15 | 5 | 23 | 5 | 6 | 5 | 47 |
| | 5 | 25 | 91878 | 25 | 27 | 25 | 23 | 25 | 39 | 25 | 46 |
| | 6 | 323 | 468 | 323 | 106 | 448 | 178 | 348 | 2685 | 323 | 187 |
| | 7 | 1456 | 37212 | 224 | 81 | 280 | 75 | 264 | 852 | 224 | 121 |
| | 8 | 20 | 21 | 20 | 23 | 20 | 15 | 20 | 61 | 20 | 9 |
| | 9 | 4200 | 30506 | 4200 | 739 | 4200 | 21181 | 4200 | 366673 | 4200 | 16875 |

TABLE : Results of the execution to obtain UCQ (time in ms)

CONCLUSIONS I

A benchmark should consider how the input represents reality wrt:

■ queries

- syntax (SELECT, COUNT, MAX, ...)
- expressiveness (CQ, UCQ, comparisons, arithmetic operations, ...)
- shape (star shaped, linear, cyclic, ...)
- size (number of atoms, number of clauses, triples...)

■ ontologies

- expressiveness (from DL-Lite_g to...)
- shape (flat, hierarchical, cyclic ...)
- size (number of concepts, properties, individuals, ...)

■ Additional information

- mappings
- ABox dependencies / EBox
- caching for several queries

CONCLUSIONS I

A benchmark should consider how the input represents reality wrt:

■ queries

- syntax (SELECT, COUNT, MAX, ...)
- expressiveness (CQ, UCQ, comparisons, arithmetic operations, ...)
- shape (star shaped, linear, cyclic, ...)
- size (number of atoms, number of clauses, triples...)

■ ontologies

- expressiveness (from DL-Lite_R to...)
- shape (flat, hierarchical, cyclic ...)
- size (number of concepts, properties, individuals, ...)

■ Additional information

- mappings
- ABox dependencies / EBox
- caching for several queries

CONCLUSIONS I

A benchmark should consider how the input represents reality wrt:

- queries

- syntax (SELECT, COUNT, MAX, ...)
- expressiveness (CQ, UCQ, comparisons, arithmetic operations, ...)
- shape (star shaped, linear, cyclic, ...)
- size (number of atoms, number of clauses, triples...)

- ontologies

- expressiveness (from DL-Lite_R to...)
- shape (flat, hierarchical, cyclic ...)
- size (number of concepts, properties, individuals, ...)

- Additional information

- mappings
- ABox dependencies / EBox
- caching for several queries

CONCLUSIONS I

A benchmark should consider how the input represents reality wrt:

■ queries

- syntax (SELECT, COUNT, MAX, ...)
- expressiveness (CQ, UCQ, comparisons, arithmetic operations, ...)
- shape (star shaped, linear, cyclic, ...)
- size (number of atoms, number of clauses, triples...)

■ ontologies

- expressiveness (from DL-Lite_R to...)
- shape (flat, hierarchical, cyclic ...)
- size (number of concepts, properties, individuals, ...)

■ Additional information

- mappings
- ABox dependencies / EBox
- caching for several queries

CONCLUSIONS I

A benchmark should consider how the input represents reality wrt:

■ queries

- syntax (SELECT, COUNT, MAX, ...)
- expressiveness (CQ, UCQ, comparisons, arithmetic operations, ...)
- shape (star shaped, linear, cyclic, ...)
- size (number of atoms, number of clauses, triples...)

■ ontologies

- expressiveness (from DL-Lite_R to...)
- shape (flat, hierarchical, cyclic ...)
- size (number of concepts, properties, individuals, ...)

■ Additional information

- mappings
- ABox dependencies / EBox
- caching for several queries

CONCLUSIONS I

A benchmark should consider how the input represents reality wrt:

- queries

- syntax (SELECT, COUNT, MAX, ...)
- expressiveness (CQ, UCQ, comparisons, arithmetic operations, ...)
- shape (star shaped, linear, cyclic, ...)
- size (number of atoms, number of clauses, triples...)

- ontologies

- expressiveness (from DL-Lite_R to...)
- shape (flat, hierarchical, cyclic ...)
- size (number of concepts, properties, individuals, ...)

- Additional information

- mappings
- ABox dependencies / EBox
- caching for several queries

CONCLUSIONS I

A benchmark should consider how the input represents reality wrt:

■ queries

- syntax (SELECT, COUNT, MAX, ...)
- expressiveness (CQ, UCQ, comparisons, arithmetic operations, ...)
- shape (star shaped, linear, cyclic, ...)
- size (number of atoms, number of clauses, triples...)

■ ontologies

- expressiveness (from DL-Lite_R to...)
- shape (flat, hierarchical, cyclic ...)
- size (number of concepts, properties, individuals, ...)

■ Additional information

- mappings
- ABox dependencies / EBox
- caching for several queries

CONCLUSIONS I

A benchmark should consider how the input represents reality wrt:

- queries

- syntax (SELECT, COUNT, MAX, ...)
- expressiveness (CQ, UCQ, comparisons, arithmetic operations, ...)
- shape (star shaped, linear, cyclic, ...)
- size (number of atoms, number of clauses, triples...)

- ontologies

- expressiveness (from DL-Lite_R to...)
- shape (flat, hierarchical, cyclic ...)
- size (number of concepts, properties, individuals, ...)

- Additional information

- mappings
- ABox dependencies / EBox
- caching for several queries

CONCLUSIONS I

A benchmark should consider how the input represents reality wrt:

- queries

- syntax (SELECT, COUNT, MAX, ...)
- expressiveness (CQ, UCQ, comparisons, arithmetic operations, ...)
- shape (star shaped, linear, cyclic, ...)
- size (number of atoms, number of clauses, triples...)

- ontologies

- expressiveness (from DL-Lite_R to...)
- shape (flat, hierarchical, cyclic ...)
- size (number of concepts, properties, individuals, ...)

- Additional information

- mappings
- ABox dependencies / EBox
- caching for several queries

CONCLUSIONS I

A benchmark should consider how the input represents reality wrt:

- queries

- syntax (SELECT, COUNT, MAX, ...)
- expressiveness (CQ, UCQ, comparisons, arithmetic operations, ...)
- shape (star shaped, linear, cyclic, ...)
- size (number of atoms, number of clauses, triples...)

- ontologies

- expressiveness (from DL-Lite_R to...)
- shape (flat, hierarchical, cyclic ...)
- size (number of concepts, properties, individuals, ...)

- Additional information

- mappings
- ABox dependencies / EBox
- caching for several queries

CONCLUSIONS I

A benchmark should consider how the input represents reality wrt:

- queries

- syntax (SELECT, COUNT, MAX, ...)
- expressiveness (CQ, UCQ, comparisons, arithmetic operations, ...)
- shape (star shaped, linear, cyclic, ...)
- size (number of atoms, number of clauses, triples...)

- ontologies

- expressiveness (from DL-Lite_R to...)
- shape (flat, hierarchical, cyclic ...)
- size (number of concepts, properties, individuals, ...)

- Additional information

- mappings
- ABox dependencies / EBox
- caching for several queries

CONCLUSIONS I



A benchmark should consider how the input represents reality wrt:

- queries

- syntax (SELECT, COUNT, MAX, ...)
- expressiveness (CQ, UCQ, comparisons, arithmetic operations, ...)
- shape (star shaped, linear, cyclic, ...)
- size (number of atoms, number of clauses, triples...)

- ontologies

- expressiveness (from DL-Lite_R to...)
- shape (flat, hierarchical, cyclic ...)
- size (number of concepts, properties, individuals, ...)

- Additional information

- mappings
- ABox dependencies / EBox
- caching for several queries

CONCLUSIONS I

A benchmark should consider how the input represents reality wrt:

- queries
 - syntax (SELECT, COUNT, MAX, ...)
 - expressiveness (CQ, UCQ, comparisons, arithmetic operations, ...)
 - shape (star shaped, linear, cyclic, ...)
 - size (number of atoms, number of clauses, triples...)
- ontologies
 - expressiveness (from DL-Lite_R to...)
 - shape (flat, hierarchical, cyclic ...)
 - size (number of concepts, properties, individuals, ...)
- Additional information
 - mappings
 - ABox dependencies / EBox
 - caching for several queries

CONCLUSIONS II

A benchmark should consider what the output means wrt:

- Shape of rewritten queries

- expressiveness
- types of clauses
- syntax with special characteristics

- Size of rewritten queries

- number of clauses
- number of atoms (and distinct atoms)
- number of joins

CONCLUSIONS II

A benchmark should consider what the output means wrt:

- Shape of rewritten queries
 - expressiveness
 - types of clauses
 - syntax with special characteristics
- Size of rewritten queries
 - number of clauses
 - number of atoms (and distinct atoms)
 - number of joins

CONCLUSIONS II

A benchmark should consider what the output means wrt:

- Shape of rewritten queries
 - expressiveness
 - types of clauses
 - syntax with special characteristics
- Size of rewritten queries
 - number of clauses
 - number of atoms (and distinct atoms)
 - number of joins

CONCLUSIONS II

A benchmark should consider what the output means wrt:

- Shape of rewritten queries
 - expressiveness
 - types of clauses
 - syntax with special characteristics
- Size of rewritten queries
 - number of clauses
 - number of atoms (and distinct atoms)
 - number of joins

CONCLUSIONS II

A benchmark should consider what the output means wrt:

- Shape of rewritten queries
 - expressiveness
 - types of clauses
 - syntax with special characteristics
- Size of rewritten queries
 - number of clauses
 - number of atoms (and distinct atoms)
 - number of joins

CONCLUSIONS II

A benchmark should consider what the output means wrt:

- Shape of rewritten queries
 - expressiveness
 - types of clauses
 - syntax with special characteristics
- Size of rewritten queries
 - number of clauses
 - number of atoms (and distinct atoms)
 - number of joins

CONCLUSIONS II

A benchmark should consider what the output means wrt:

- Shape of rewritten queries
 - expressiveness
 - types of clauses
 - syntax with special characteristics
- Size of rewritten queries
 - number of clauses
 - number of atoms (and distinct atoms)
 - number of joins

CONCLUSIONS II

A benchmark should consider what the output means wrt:

- Shape of rewritten queries
 - expressiveness
 - types of clauses
 - syntax with special characteristics
- Size of rewritten queries
 - number of clauses
 - number of atoms (and distinct atoms)
 - number of joins

REFERENCES I

- Diego Calvanese, Giuseppe De Giacomo, Domenico Lembo, Maurizio Lenzerini, and Riccardo Rosati. Tractable reasoning and efficient query answering in description logics: The DL-Lite family. *Journal of Automated Reasoning*, 39(3):385–429, October 2007. doi: 10.1007/s10817-007-9078-x. URL <http://dx.doi.org/10.1007/s10817-007-9078-x>.
- Alexandros Chortaras, Despoina Trivela, and Giorgos Stamou. Optimized query rewriting for OWL 2 QL. In Nikolaj Björner and Viorica Sofronie-Stokkermans, editors, *Automated Deduction – CADE-23*, volume 6803, pages 192–206. Springer Berlin Heidelberg, Berlin, Heidelberg, 2011. ISBN 978-3-642-22437-9. URL <http://www.springerlink.com/content/g8153m783k638210/>.
- Thomas Eiter, Magdalena Ortiz, Mantas Vsimkus, Trung-Kien Tran, and Guohui Xiao. Query rewriting for horn-SHIQ plus rules. In *Proc. of the 26th AAAI Conference on Artificial Intelligence, AAAI, 2012*. URL <http://www.aaai.org/ocs/index.php/AAAI/AAAI12/paper/viewPDFInterstitial/4931/5263>.
- Georg Gottlob, Giorgio Orsi, and Andreas Pieris. Ontological queries: Rewriting and optimization (extended version). *arXiv:1112.0343*, December 2011. URL <http://arxiv.org/abs/1112.0343>.
- Martha Imprialou, Giorgos Stoilos, and Bernardo Cuenca Grau. Benchmarking ontology-based query rewriting systems. In *Proceedings of the Twenty-Sixth AAAI Conference on Artificial Intelligence, AAAI, 2012*. URL <http://www.aaai.org/ocs/index.php/AAAI/AAAI12/paper/viewPDFInterstitial/4910/5270>.
- Jose Mora and Oscar Corcho. Engineering optimisations in query rewriting for OBDA. In *Proceedings of the 9th International Conference on Semantic Systems, ICPS*, pages 41–48, Graz, Austria, September 2013. ACM.
- Héctor Pérez-Urbina, Ian Horrocks, and Boris Motik. Efficient query answering for OWL 2. In *The Semantic Web - ISWC 2009*, volume 5823 of *Lecture Notes in Computer Science*, pages 489–504. Springer, 2009. URL http://dx.doi.org/10.1007/978-3-642-04930-9_31.
- Riccardo Rosati. Prexto: Query rewriting under extensional constraints in DL-Lite. In Elena Simperl, Philipp Cimiano, Axel Polleres, Oscar Corcho, and Valentina Presutti, editors, *The Semantic Web: Research and Applications*, volume 7295 of *Lecture Notes in Computer Science*, pages 360–374. Springer Berlin / Heidelberg, 2012. ISBN 978-3-642-30283-1. URL <http://www.springerlink.com/content/1j61g52j78525175/abstract/>.
- Riccardo Rosati and Alessandro Almatelli. Improving query answering over DL-Lite ontologies. In Fangzhen Lin, Ulrike Sattler, and Miroslaw Truszczyński, editors, *Proceedings of the Twelfth International Conference on the Principles of Knowledge Representation and Reasoning*. AAAI Press, 2010. URL <http://dblp.uni-trier.de/db/conf/kr/kr2010.html>.
- T. Venetis, G. Stoilos, and G. Stamou. Query rewriting under query extensions for OWL 2 QL ontologies. In *The 7th International Workshop on Scalable Semantic Web Knowledge Base Systems (SSWS 2011)*, page 59, 2011. URL <http://iswc2011.semanticweb.org/fileadmin/iswc/Papers/Workshops/SSWS/SSWS2011-Proceedings.pdf>.

TOWARDS A SYSTEMATIC BENCHMARKING OF ONTOLOGY-BASED QUERY REWRITING SYSTEMS

José Mora and Óscar Corcho

{jmora, ocorcho}@fi.upm.es

Boadilla - October 10, 2013