Scalable RDF Management in the Web of Data

Toward Efficient Interchange of RDF Data Streams

Javier D. Fernández

Ontology Engineering Group
Departamento de Inteligencia Artificial
Facultad de Informática
Universidad Politécnica de Madrid

Date: 24/04/2014

Index

1. Background
2. Motivation
3. What have I done so far?
4. What am I currently doing?
5. Next steps

What is Linked Data?

1. Use URIs as names for things.

2. Use HTTP URIs so that people can look up those names.

3. When someone looks up a URI, provide useful information.

4. Include links to other URIs.

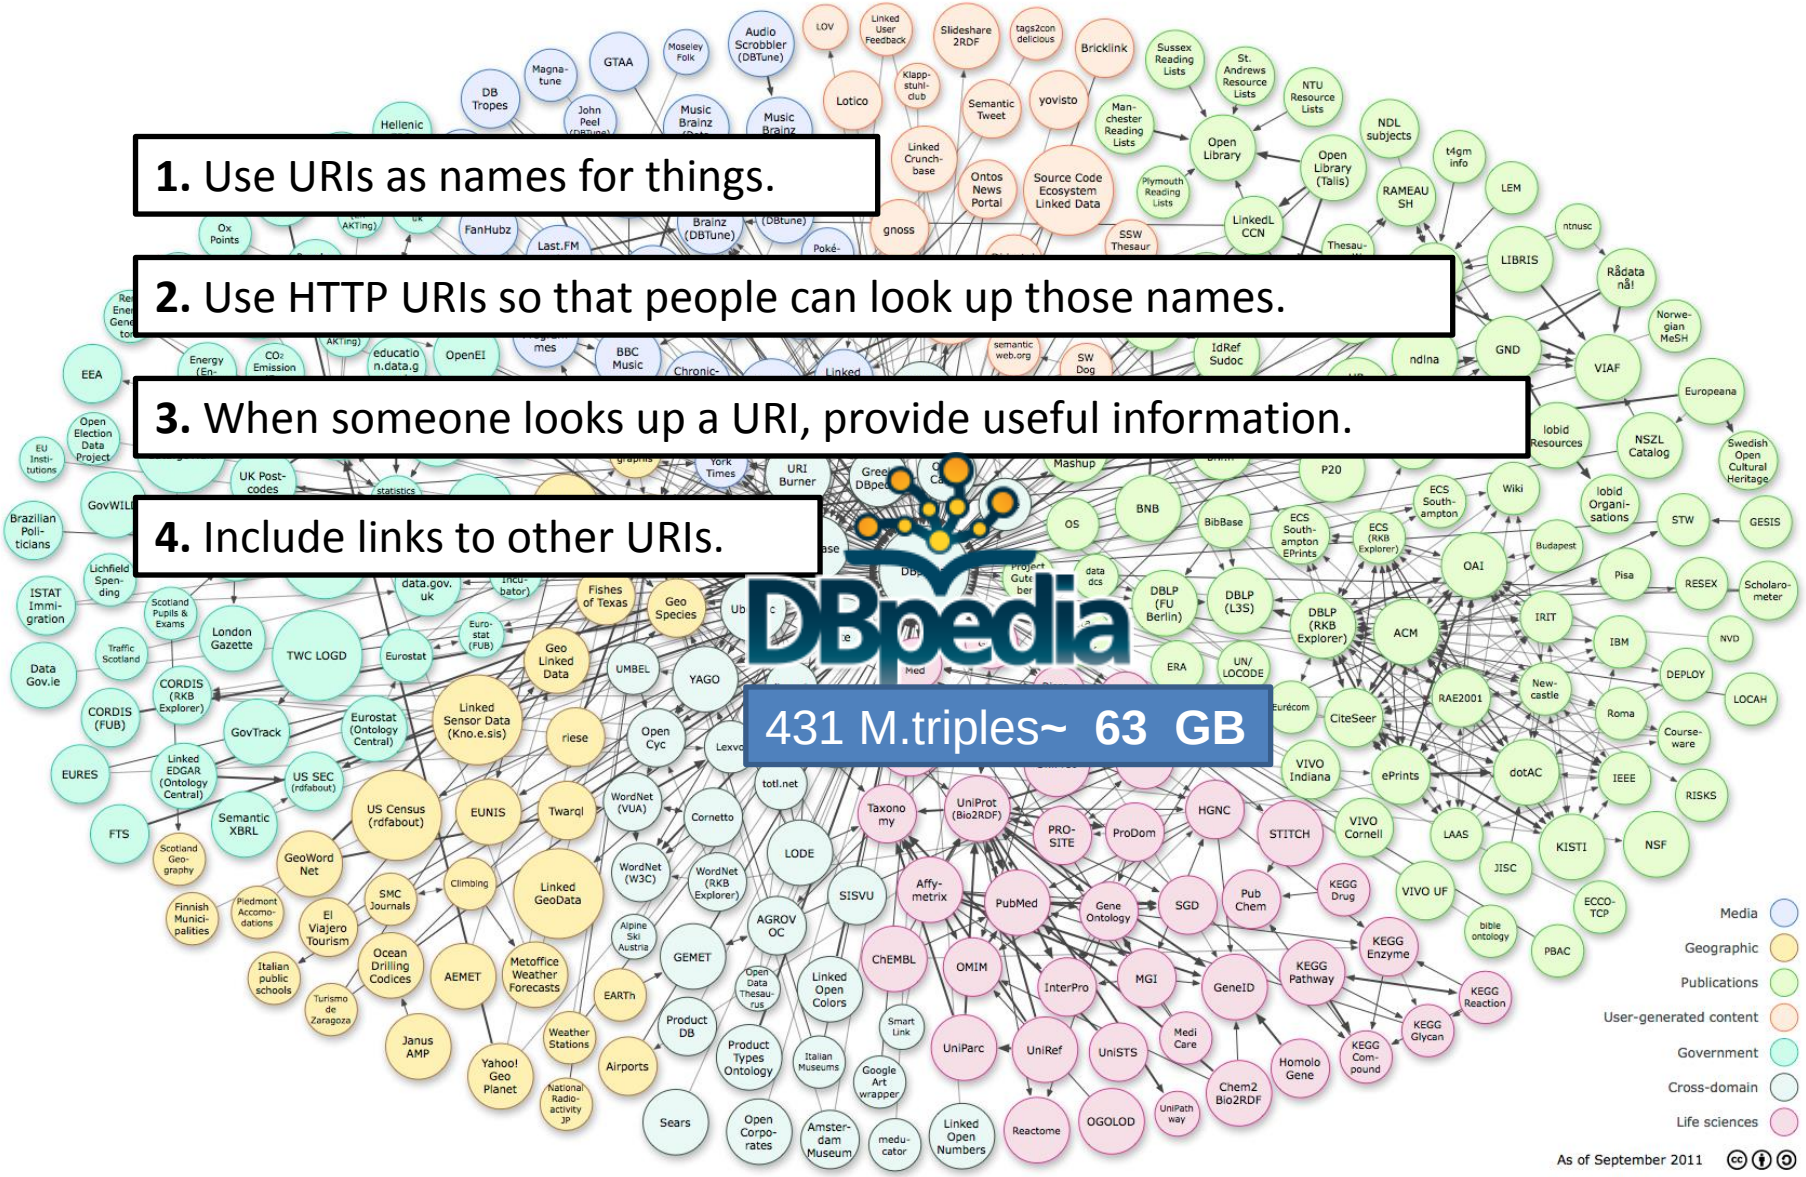

Linked Data cloud: > 62 billion triples

What is Linked Data?

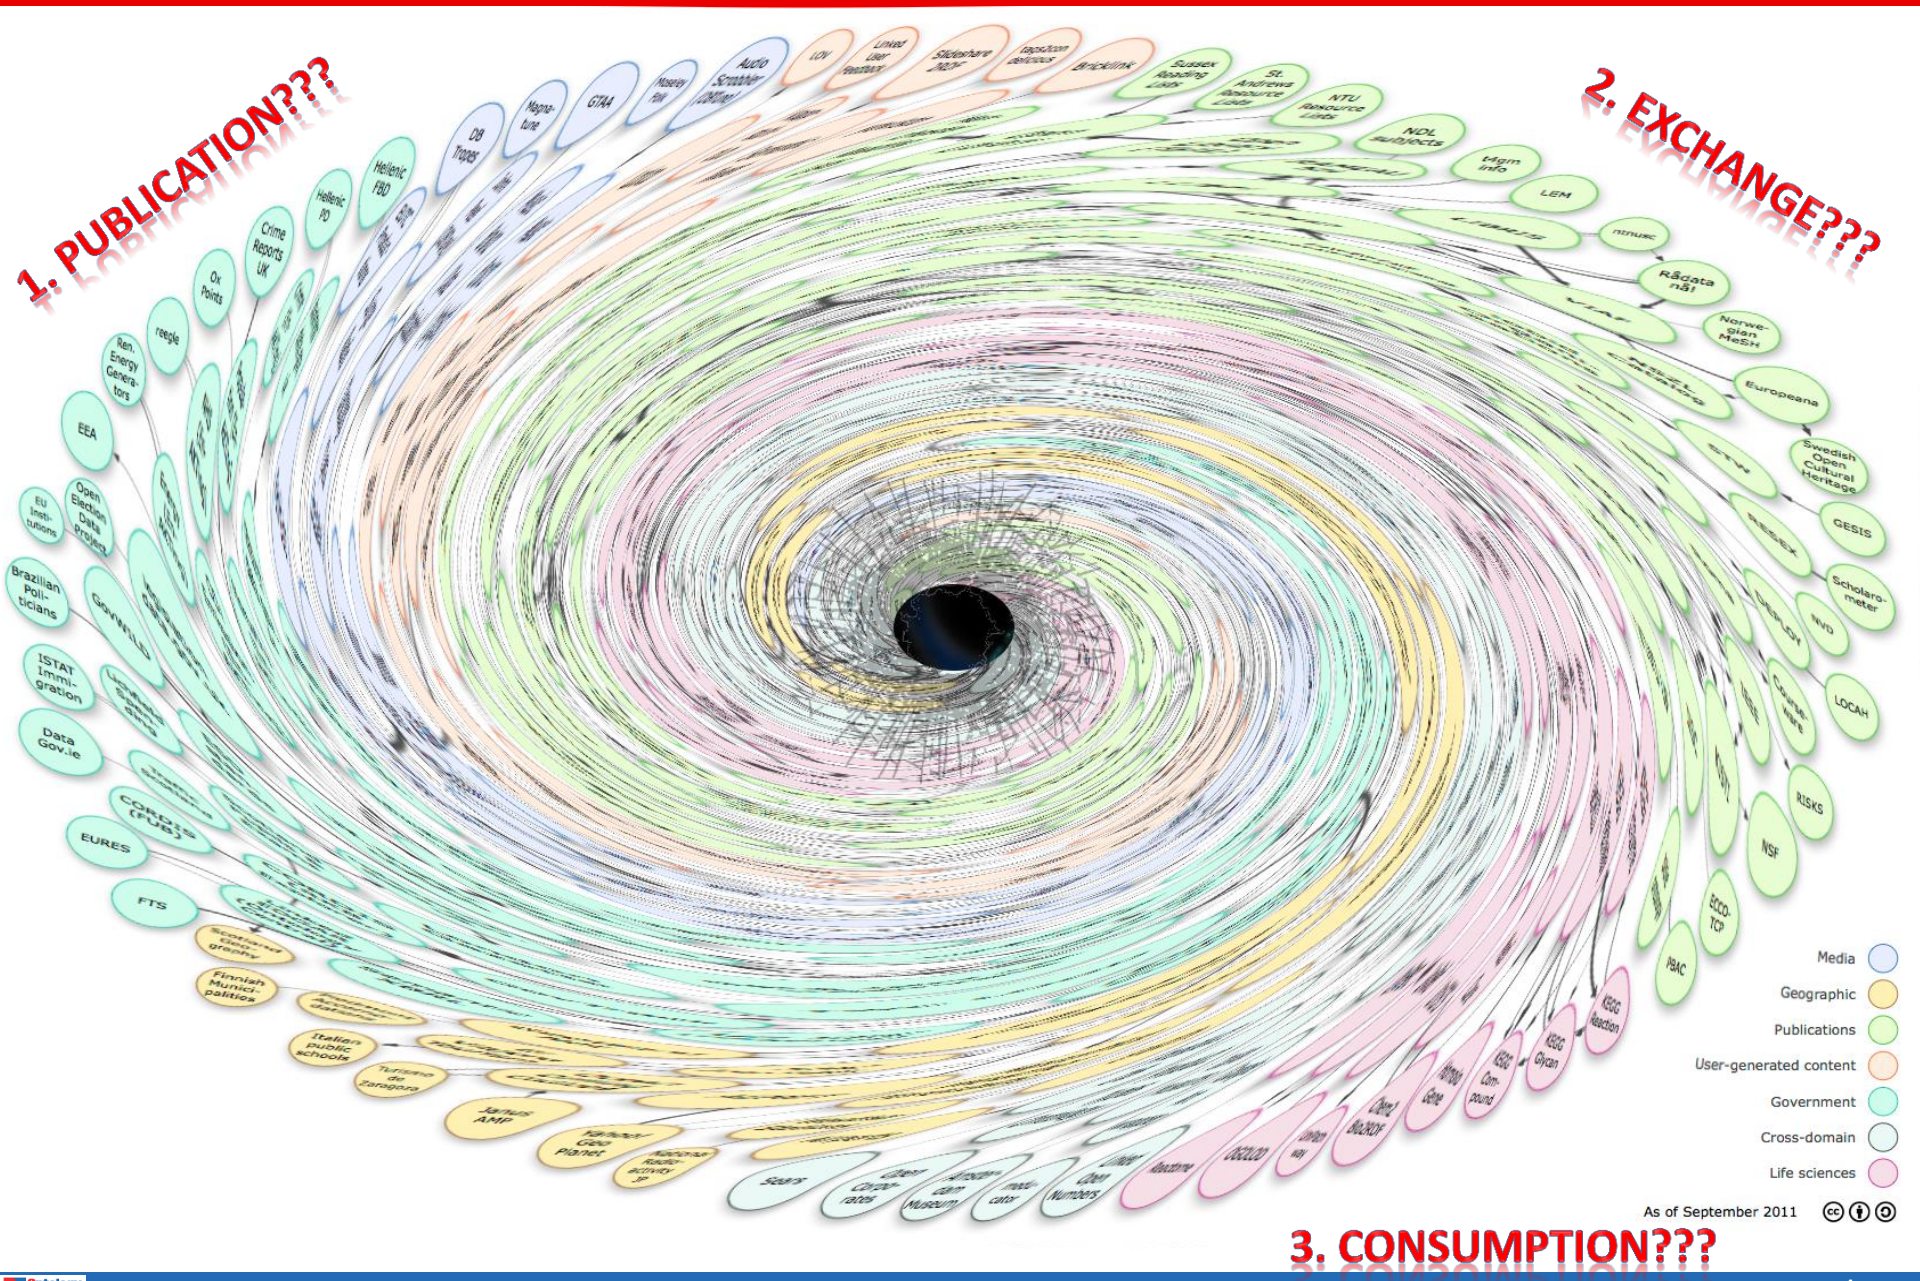

Publication & Exchange: Is this the best option?

Dataset	en	ca	de
Mapping-based Types	nt ?	nt ?	nt ?
	nq ?	nq ?	nq ?
	ttl ?	ttl ?	ttl ?
Mapping-based Types (Heuristic)	nt ?	--	--
	--	--	--
	ttl ?	--	--
Mapping-based Properties	nt ?	nt ?	nt ?
	nq ?	nq ?	nq ?
	ttl ?	ttl ?	ttl ?
Mapping-based Properties (Cleaned)	nt ?	--	--
	nq ?	--	--
	ttl ?	--	--
Mapping-based Properties (Specific)	nt ?	nt ?	nt ?
	nq ?	nq ?	nq ?
	ttl ?	ttl ?	ttl ?

Dataset	en	ca	de	es	eu
Titles	nt ?	nt ?	nt ?	nt ?	nt ?
	nq ?	nq ?	nq ?	nq ?	nq ?
	ttl ?	ttl ?	ttl ?	ttl ?	ttl ?

DBpedia

431 M.triples~ 63 GB

Publication & Exchange: Is this the best option?

Dataset	en	ca	de
Mapping-based Types	nt ?	nt ?	nt ?
	nq ?	nq ?	nq ?
Mapping-based Properties (Cleaned)	nq ?	--	--
	ttl ?	--	--
Mapping-based Properties (Specific)	nt ?	nt ?	nt ?
	nq ?	nq ?	nq ?
	ttl ?	ttl ?	ttl ?

Dataset	en	ca	de	es	eu
Titles	nt ?	nt ?	nt ?	nt ?	nt ?
	nq ?	nq ?	nq ?	nq ?	nq ?
	ttl ?	ttl ?	ttl ?	ttl ?	ttl ?

"the published RDF dumps are actually bulks with no structure, no design, no final user in mind. They resemble unwanted creatures whose owners are keen to be rid of them"

Claudio Gutiérrez

Drawbacks of Current RDF Serialization

- Very verbose
  - Designed for human readability (not for machines)
  - HUGE → text compression/decompression
- Lack of (standard) metadata
  - "What is this?" phenomenon
- Search offline
  - Scan the whole exchanged dump.
  - (decompress)+ index the file + search

- Very verbose
- for machines)
- Lac
  - "Wh
- dex the file

In this RDF deluge,
if RDF is meant to be machine processable,
why are we using plain RDF??

Given an RDF dataset, potentially huge, a lightweight binary RDF can encode the data leveraging the skewed structure of RDF graphs for the purposes of

1. Large spatial savings,
2. easy and modular data-centric publication and parsing and
3. data retrieval.

Applications:

- Publish a large dataset on the Web.
- Transfer between two servers.
- Distributed RDF Data Management.
- Fast In-Memory Query Engine.

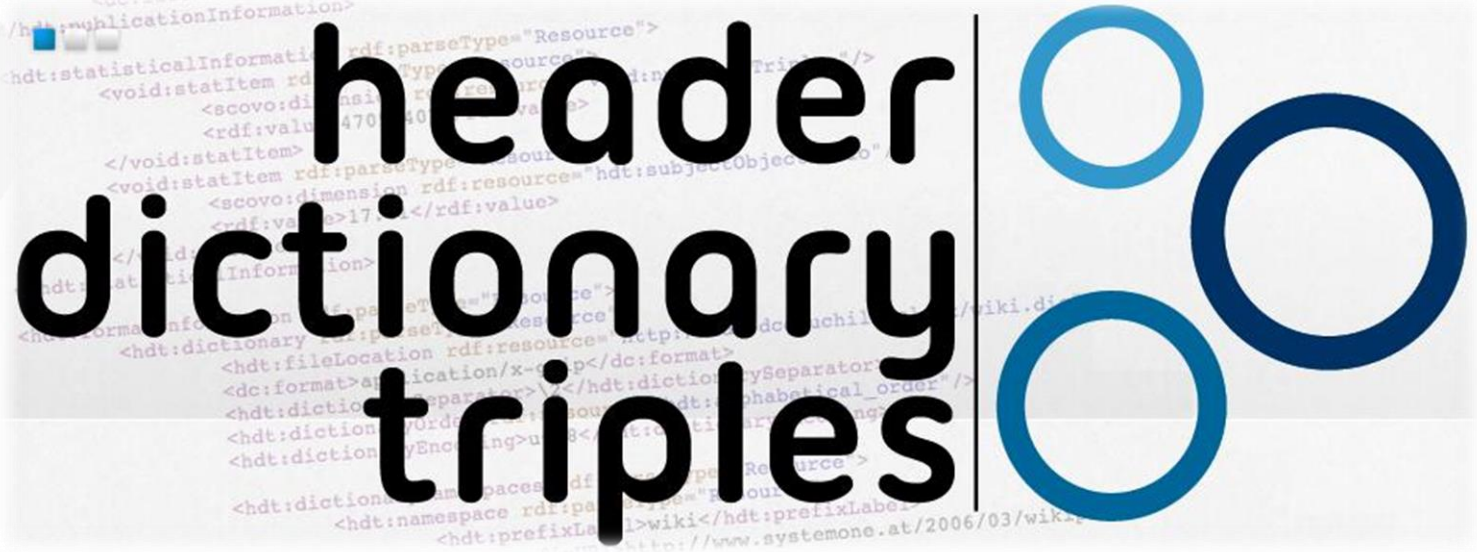

- Binary Serialization of RDF
- Highly compact
- Includes indexes to solve SPARQL Triple Patterns once it is loaded in main memory
- W3C Submission. http://www.w3.org/Submission/2011/03/

HDT (Header-Dictionary-Triples) Overview

RDF

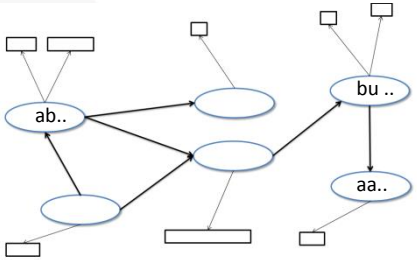

hdt

Header

◈metadata describing the RDF dataset

Dictionary

1 aa..
2 ab..
3 bu
...
...

◈Mapping between IDs ←→elements in the dataset

Triples

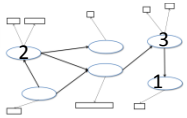

◈Structure of the data after the ID replacement

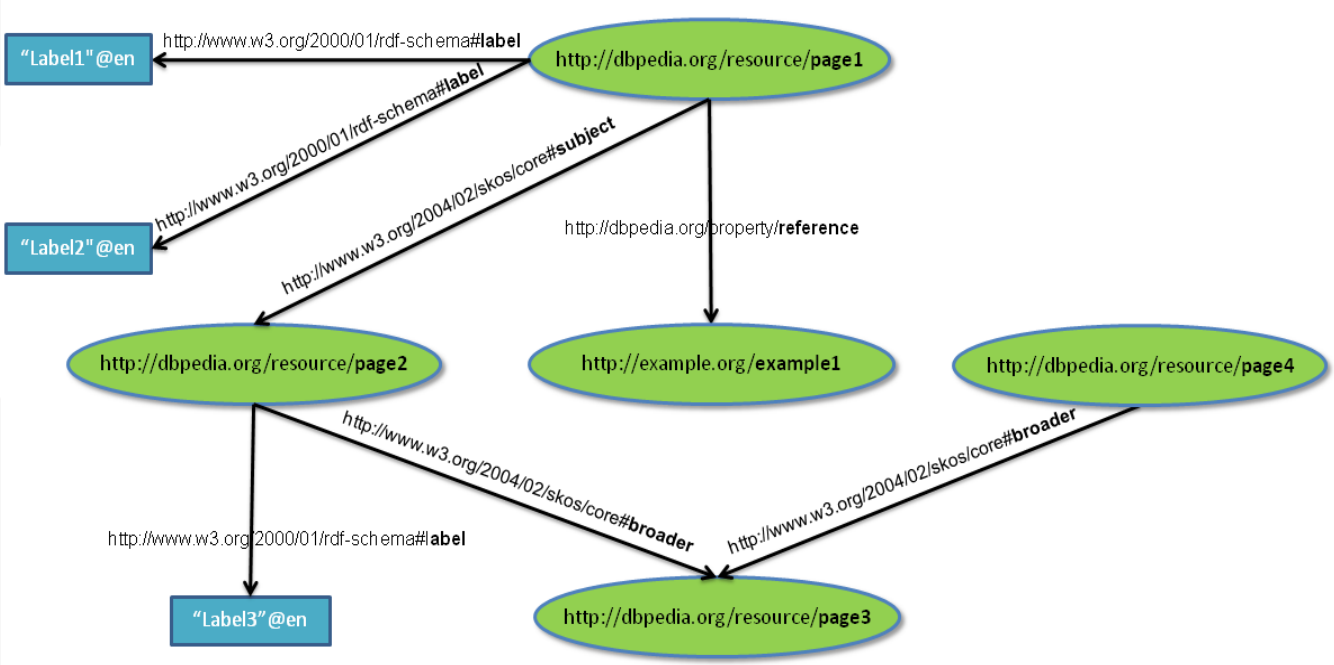

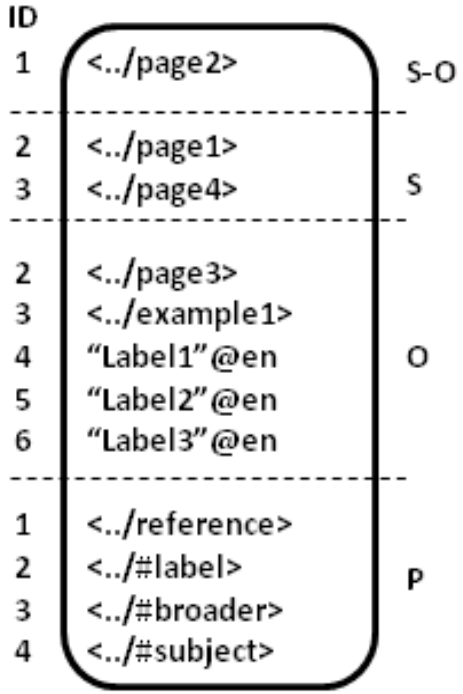

- Mapping of strings to correlative IDs. {1..n}
- Lexicographically sorted, no duplicates.
- Front Coding for each section.

Triples

1 2 6
1 3 2
2 1 3
2 2 4
2 2 5
2 4 1
3 3 2

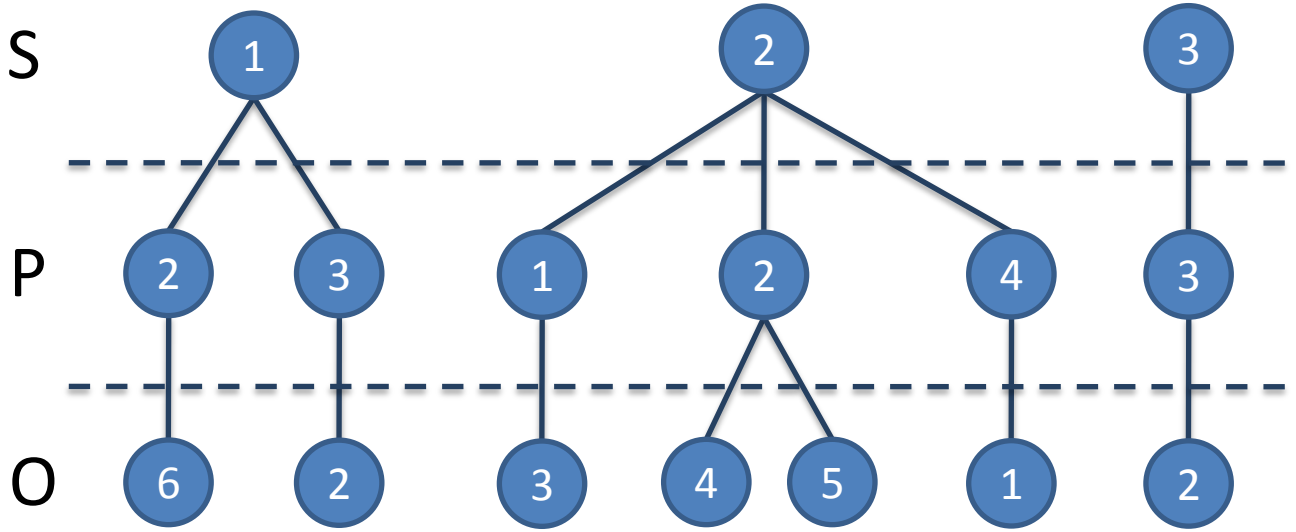

Array Y	2	3	1	2	4	3	
Bitmap Y	1	0	1	0	0	1	
Array Z	6	2	3	4	5	1	2
Bitmap Z	1	1	1	1	0	1	1


DEMO
Time

FOUR current ways of consuming HDT

rdfhdt.org

- [1] Command line Tool:
  - Export/Import
    - $ rdf2hdt file.nt output.hdt
    - $ hdt2rdf file.hdt output.nt
  - Query
    - $ hdtsearch file.hdt

FOUR current ways of consuming HDT

- [2] C++/Java Library -> Use within Jena!

```java
// Load HDT file
QueryableHDT hdt = HDTFactory.createQueryableHDT();
hdt.loadFromHDT("data/example.hdt", null);
hdt.loadOrCreateIndex(null);

// Search pattern: Empty string means "any"
IteratorTripleString it = hdt.search("", "", "");
while(it.hasNext()) {
    TripleString ts = it.next();
    System.out.println(ts);
}

// Create Jena Model on top of HDT.
HDTGraph graph = new HDTGraph(hdt);
Model model = new ModelCom(graph);
```

FOUR current ways of consuming HDT

- [3] Web Service:
  - Import into HDT
    - http://srvgal85.deri.ie/hdt-online/
    - Thanks to Michael Hausenblas.

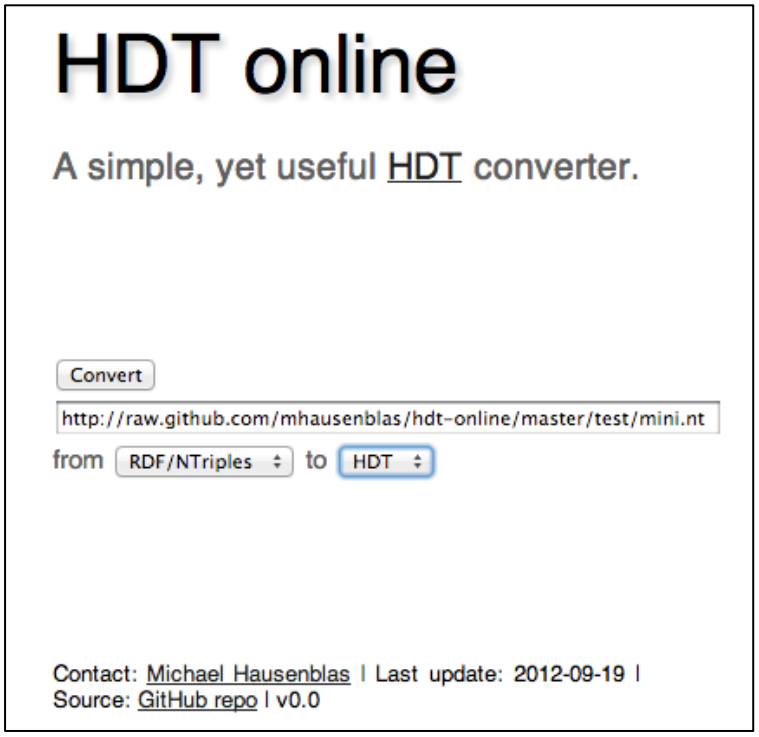

FOUR current ways of consuming HDT

- [4] Desktop tool HDT-it!
  - Thanks to Mario Arias (DERI)

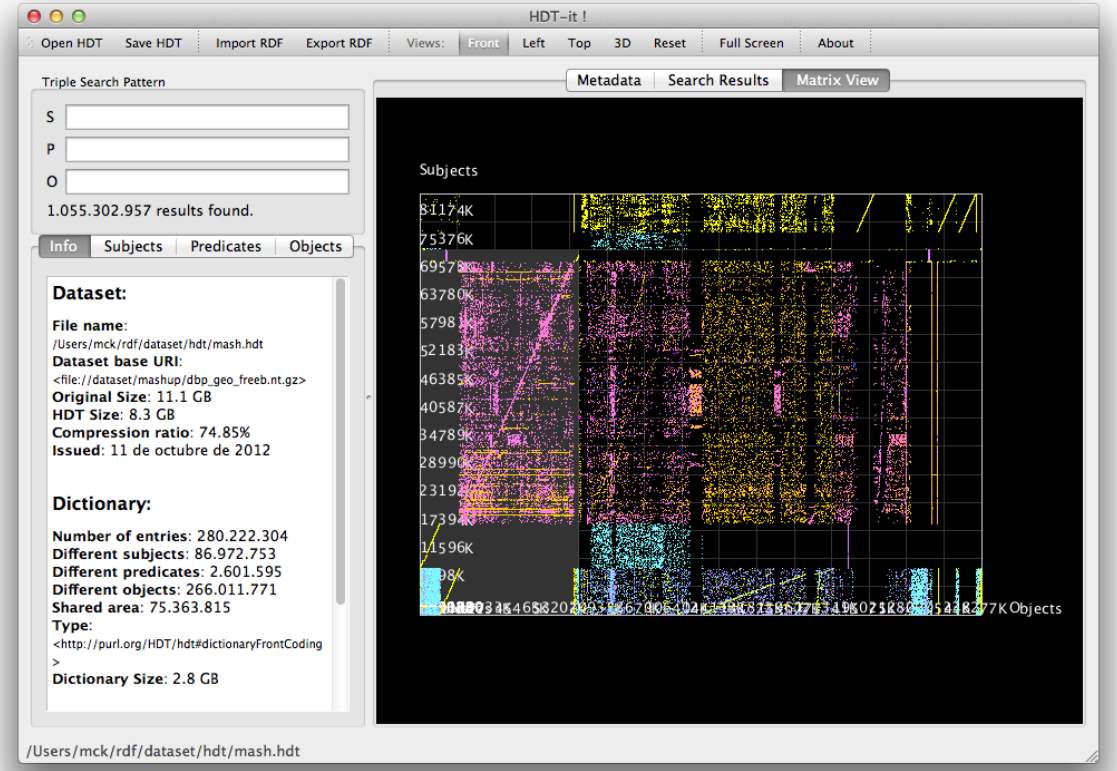

- Data is ready to be consumed 10-15x faster.
  - Exchange time reduced.
  - Indexing burden on server = Lightweight client processing.
- Competitive query performance.
  - Very fast on triple patterns.
  - Joins on the same scale of existing solutions.
- This is useful for applications that…
  - need a fast, compact read-only in-memory RDF store.
  - consider a static view of RDF datasets
  - want to share self-queryable RDF dumps.
  - need fast download & query.

Toward Efficient Interchange of RDF

Data Streams

Running example. Preliminary hypothesis & Research Questions

- 468 stations
- 4.3 M users/day

PH1.- Given a set of RDF data streams, it is possible to define an RDF interchange format that optimizes the space and time for the data exchange and parsing.

PH2.- Given an RDF streaming engine, and a set of SPARQL queries, a RDF interchange format can be tuned to offer better performance in data exchange among processing nodes and query resolution.

RQ1.- Is HDT a good solution for these dynamic data? NO
RQ2.- Which are the particularities of RDF data streams?
RQ3.- Can an RDF interchange format be parallelizable for compression and decompression (parsing)

sens-obs:Observation_AirTemperature_4UT01_2003_3_31_22_15_00

a weather:TemperatureObservation ;
om-owl:observedProperty weather:_AirTemperature ;
om-owl:procedure sens-obs:System_4UT01 ;
om-owl:result sens-obs:MeasureData_AirTemperature_4UT01_2003_3_31_22_15_00 ;
om-owl:samplingTime sens-obs:Instant_2003_3_31_22_15_00 .

sens-obs:Observation_WindGust_4UT01_2003_3_31_18_25_00

a weather:WindSpeedObservation ;
om-owl:observedProperty weather:_WindGust ;
om-owl:procedure sens-obs:System_4UT01 ;
om-owl:result sens-obs:MeasureData_WindGust_4UT01_2003_3_31_18_25_00 ;
om-owl:samplingTime sens-obs:Instant_2003_3_31_18_25_00 .

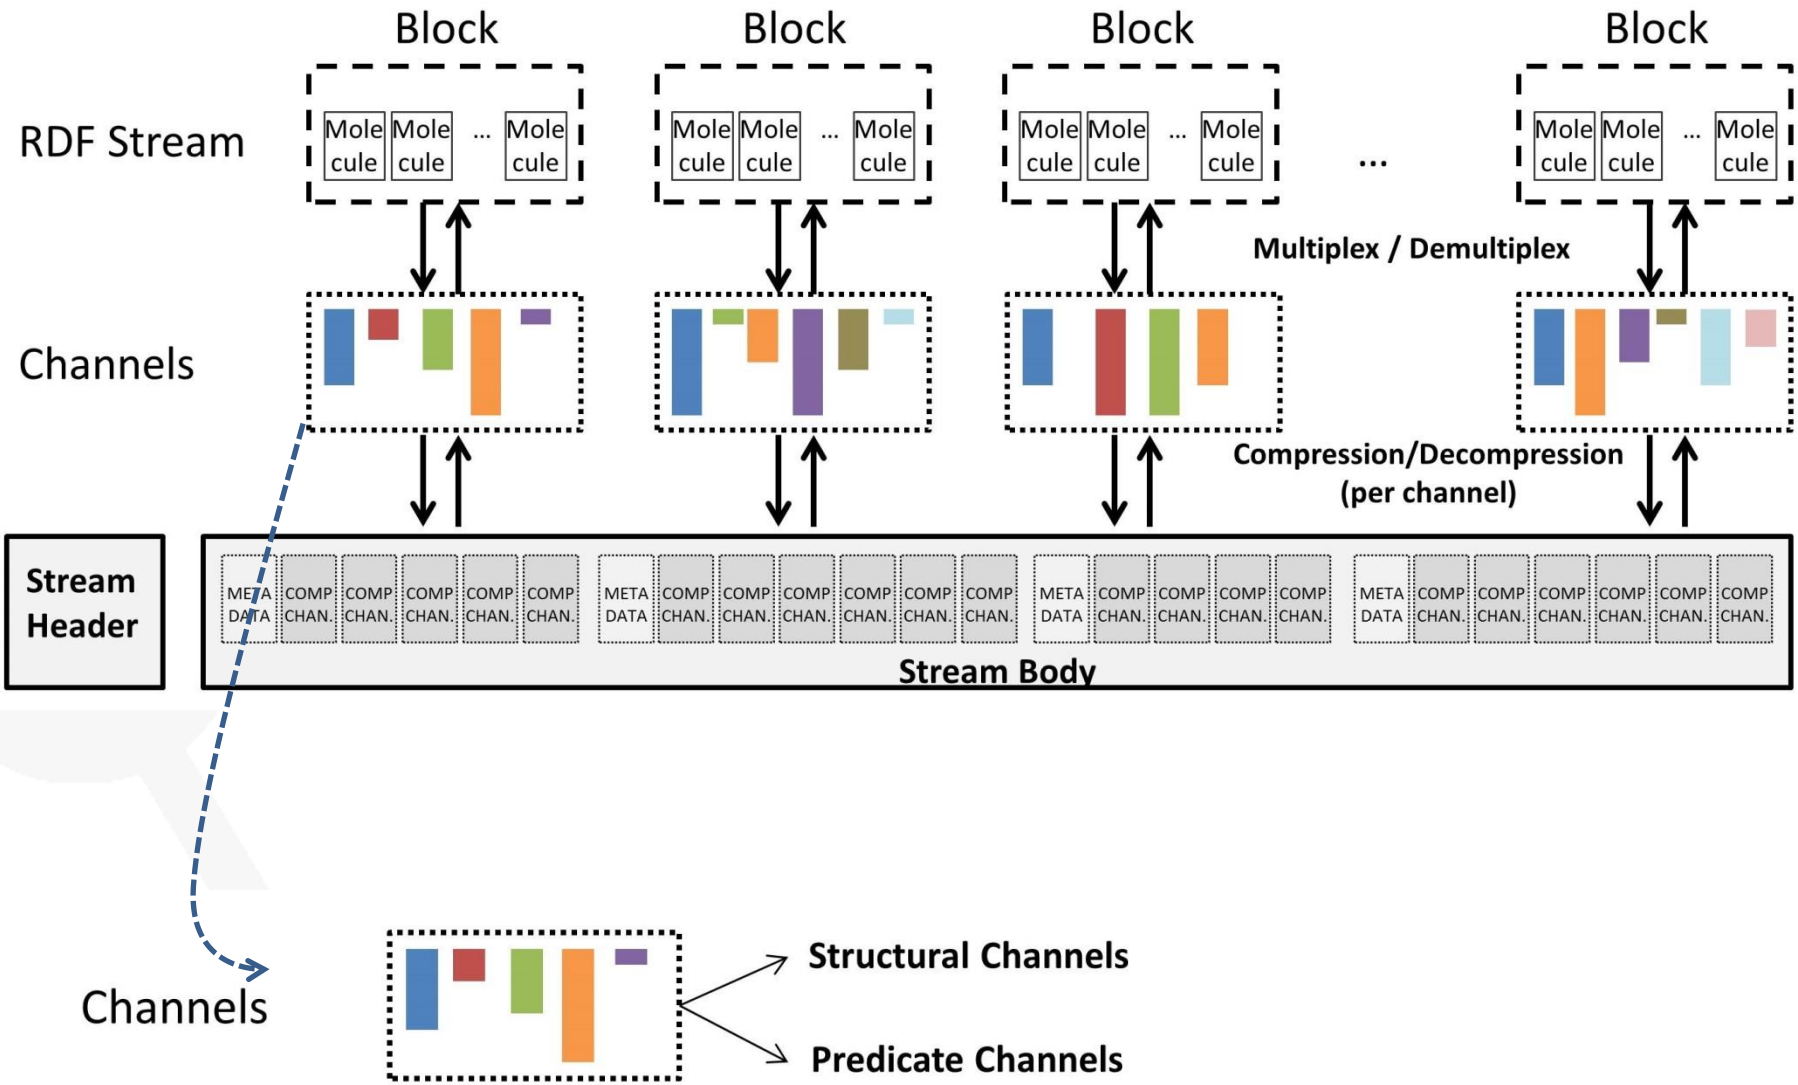

Based on: Efficient XML Interchange (EXI) format

Channels

Structural Channels

Predicate Channels

Structural Channels

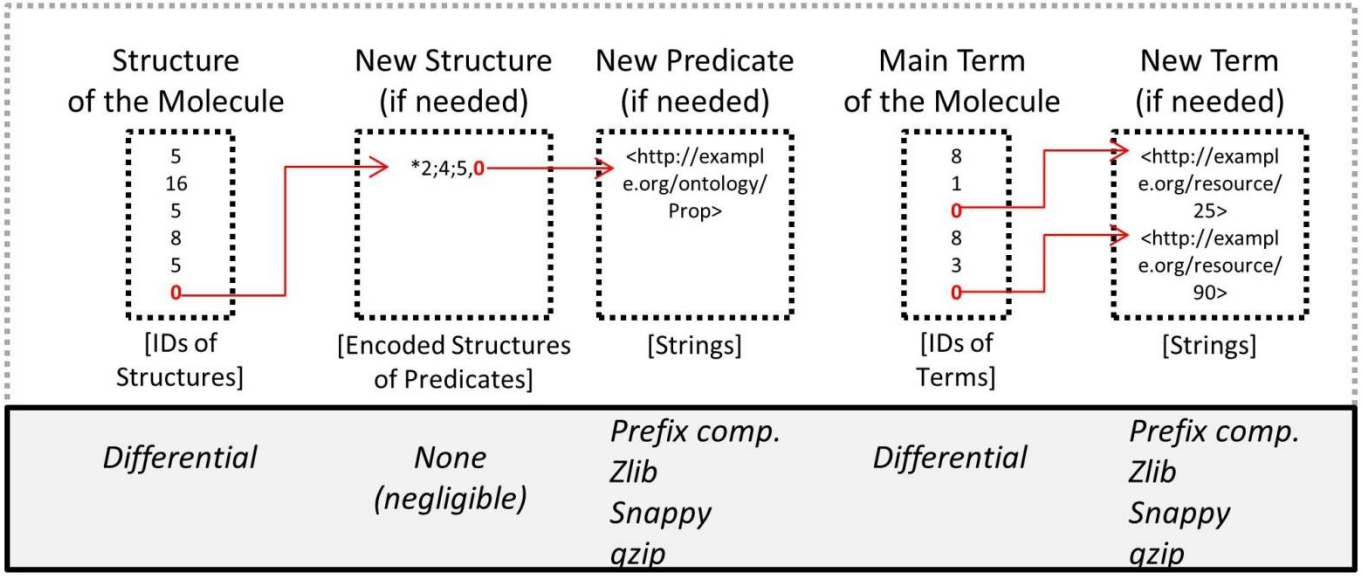

Compression possibilites

Differential	*None (negligible)*	*Prefix comp.* *Zlib* *Snappy* *gzip*	*Differential*	*Prefix comp.* *Zlib* *Snappy* *gzip*

Predicate Channels

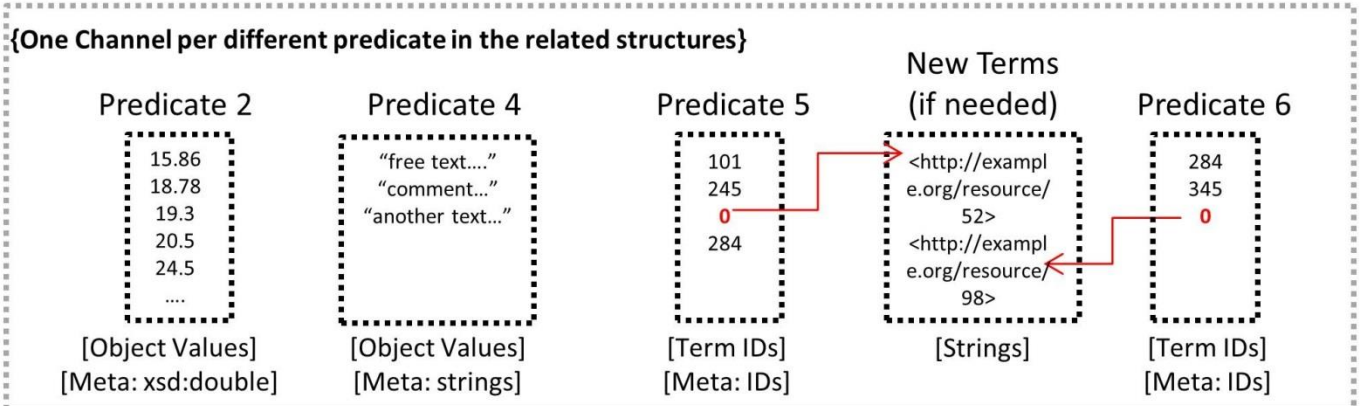

Compression possibilites

Differential	*Zlib* *Snappy* *gzip*	*Differential*	*Prefix comp.* *Zlib* *Snappy* *gzip*	*Differential*

Preliminary results Compressed Streaming

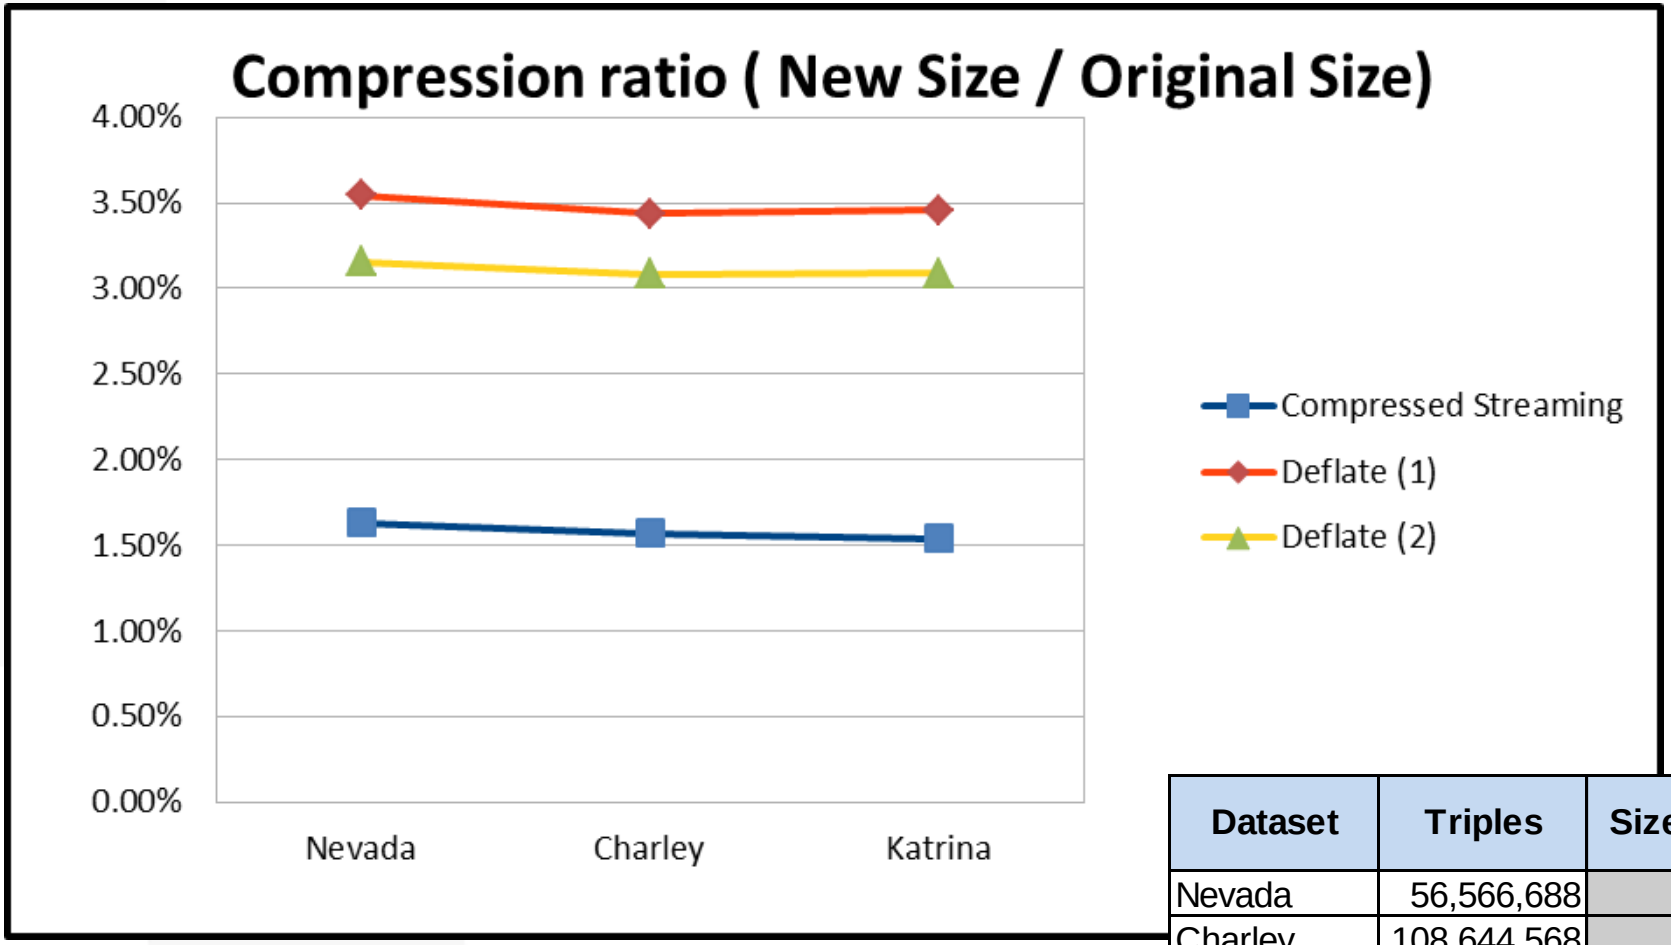

Dataset	Triples	Size (MB)
Nevada	56,566,688	7,495
Charley	108,644,568	21,470
Katrina	179,128,407	35,548

Preliminary results Compressed Streaming

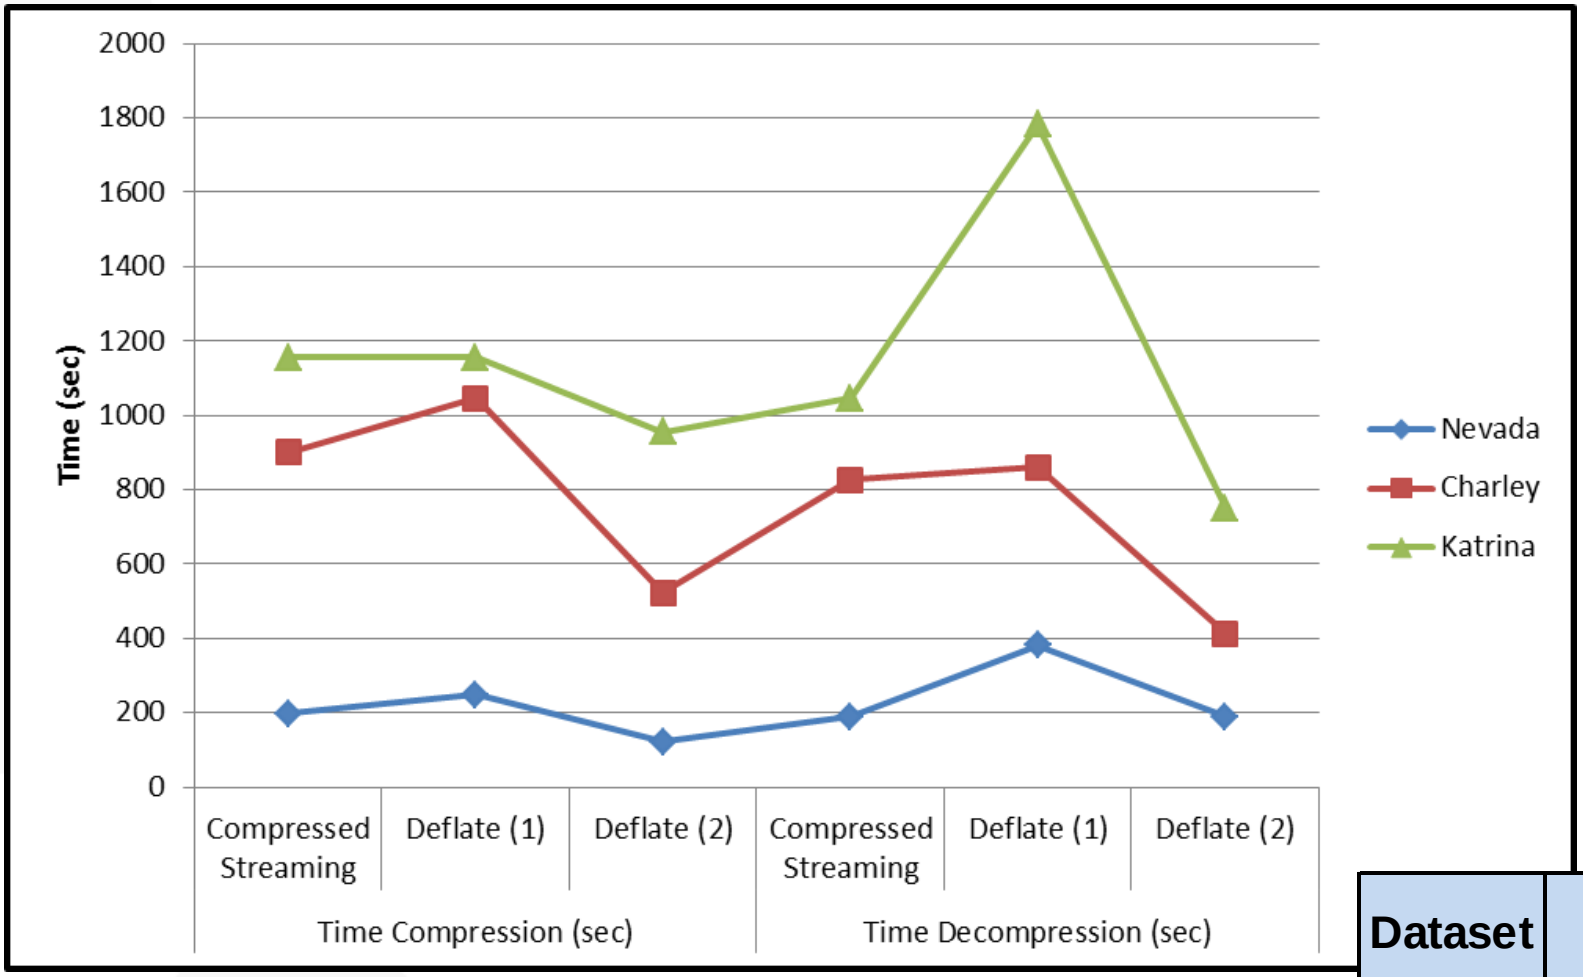

Dataset	Triples
Nevada	56,566,688
Charley	108,644,568
Katrina	179,128,407

- Finish and test the proposal with different data streams
  - Sensor and other data streams → data is welcome!
  - Release the library (Java) → feedback is welcome!
  - ISWC paper
- Parallel compression/decompression
  - preliminary proposal on Storm
- Integration within RDF streaming Engines
  - e.g. morph-streams, CQELS Cloud
  - 3 purposes:
    - scaling to higher input data rates
    - minimizing the data exchange among processing nodes
    - serving a small set of operators on the compressed data

Scalable RDF Management in the Web of Data

Toward Efficient Interchange of RDF Data Streams

Javier D. Fernández

Ontology Engineering Group
Departamento de Inteligencia Artificial
Facultad de Informática
Universidad Politécnica de Madrid

Date: 24/04/2014